

Importing the libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

Data Collection & Analysis

```
In [2]: df=pd.read_csv("C:/Users/Asus/Downloads/insurance.csv")
df
```

```
Out[2]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [3]: df.shape
```

```
Out[3]: (1338, 7)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

Categorical Features:

- Sex
- Smoker
- Region

Data Analysis

```
In [5]: df.describe()
```

```
Out[5]:
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
In [6]: df.isnull().sum()
```

```
Out[6]: age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

```
In [7]: df.nunique()
```


```
Out[7]: age      47
sex        2
bmi      548
children   6
smoker     2
region     4
charges  1337
dtype: int64
```

```
In [8]: df.columns
```

```
Out[8]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')
```


Data visualization

```
In [9]: plt.figure(figsize=(6, 5))
sns.histplot(df['age'], kde=True, color='green', edgecolor='black')
plt.xlabel('Age')
plt.ylabel('Count')
plt.title('Distribution of Age')
plt.xticks(rotation=45)
plt.show()
```



```
In [10]: # Gender column using value_counts and bar plot
plt.figure(figsize=(6, 6))
df['sex'].value_counts().plot(kind='bar', color=['lightcoral', 'lightgreen'])
```

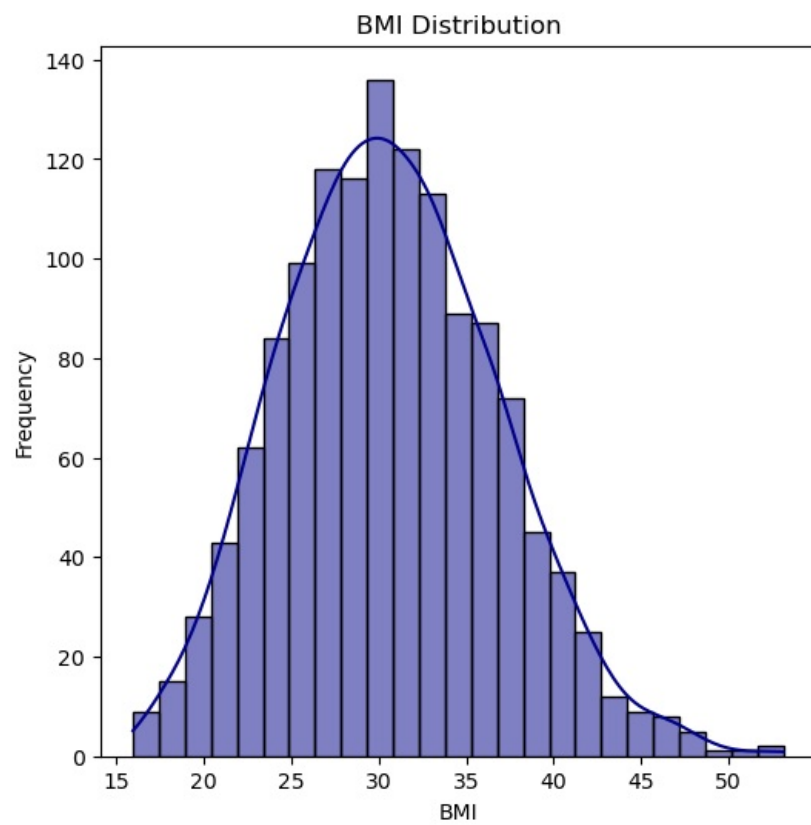
```
plt.title('Sex Distribution')
plt.xlabel('Sex')
plt.ylabel('Count')
plt.show()
```



```
In [11]: df['sex'].value_counts()
```


```
Out[11]: male      676
female    662
Name: sex, dtype: int64
```

```
In [12]: # BMI distribution using histplot with a dark color
plt.figure(figsize=(6, 6))
sns.histplot(df['bmi'], kde=True, color='darkblue')
plt.title('BMI Distribution')
plt.xlabel('BMI')
plt.ylabel('Frequency')
plt.show()
```



Normal BMI Range --> 18.5 to 24.9


```
In [14]: plt.figure(figsize=(8, 6))
colors = ['teal', 'skyblue', 'lightcoral', 'lightgreen', 'orange']
df['children'].value_counts().sort_index().plot(kind='bar', color=colors)
plt.title('Distribution of Children')
plt.xlabel('Number of Children')
plt.ylabel('Count')
plt.show()
```



```
In [15]: df['children'].value_counts()
```

```
Out[15]: 0    574
          1    324
          2    240
          3    157
          4     25
          5     18
          Name: children, dtype: int64
```


```
In [16]: # smoker column
plt.figure(figsize=(6,6))
sns.countplot(x='smoker', data=df)
plt.title('smoker')
plt.show()
```



```
In [17]: df['smoker'].value_counts()
```

```
Out[17]: no    1064
          yes     274
          Name: smoker, dtype: int64
```

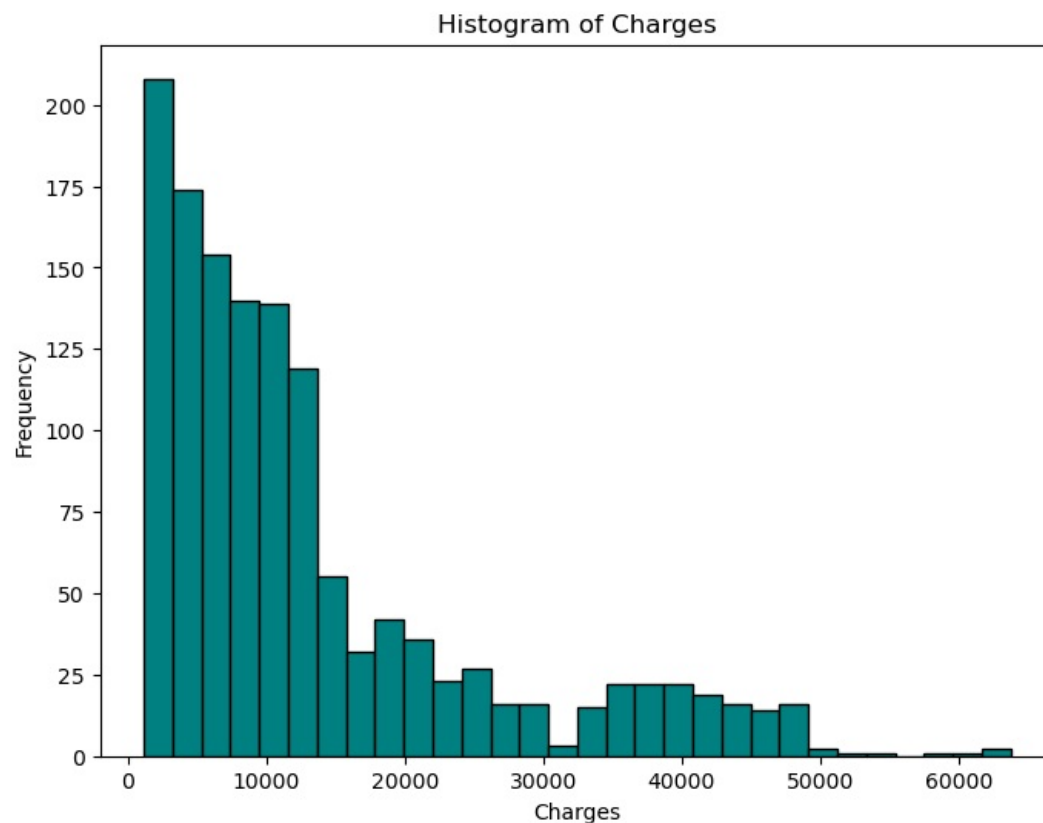
```
In [18]: # region column
plt.figure(figsize=(6,6))
sns.countplot(x='region', data=df)
plt.title('region')
plt.show()
```



```
In [19]: df['region'].value_counts()
```

```
Out[19]: southeast    364
southwest    325
northwest    325
northeast    324
Name: region, dtype: int64
```

```
In [20]: plt.figure(figsize=(8, 6))
plt.hist(df['charges'], bins=30, color='teal', edgecolor='black')
plt.title('Histogram of Charges')
plt.xlabel('Charges')
plt.ylabel('Frequency')
plt.show()
```



Data Pre-Processing

Encoding the categorical features

```
In [21]: # encoding sex column
df.replace({'sex':{'male':0,'female':1}}, inplace=True)

# encoding 'smoker' column
df.replace({'smoker':{'yes':0,'no':1}}, inplace=True)

# encoding 'region' column
df.replace({'region':{'southeast':0,'southwest':1,'northeast':2,'northwest':3}}, inplace=True)
```

Splitting the Features and Target

```
In [22]: X = df.drop(columns='charges', axis=1)
Y = df['charges']
```

```
In [23]: print(X)
```

	age	sex	bmi	children	smoker	region
0	19	1	27.900	0	0	1
1	18	0	33.770	1	1	0
2	28	0	33.000	3	1	0
3	33	0	22.705	0	1	3
4	32	0	28.880	0	1	3
...
1333	50	0	30.970	3	1	3
1334	18	1	31.920	0	1	2
1335	18	1	36.850	0	1	0
1336	21	1	25.800	0	1	1
1337	61	1	29.070	0	0	3

[1338 rows x 6 columns]

```
In [24]: print(Y)
```

0	16884.92400
1	1725.55230
2	4449.46200
3	21984.47061
4	3866.85520
...	...
1333	10600.54830
1334	2205.98080
1335	1629.83350
1336	2007.94500
1337	29141.36030

Name: charges, Length: 1338, dtype: float64

Splitting the data into Training data & Testing Data

```
In [25]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
```

```
In [26]: print(X.shape, X_train.shape, X_test.shape)
```

(1338, 6) (1070, 6) (268, 6)

Model Training

Linear Regression

```
In [27]: # loading the Linear Regression model
regressor = LinearRegression()
```

```
In [28]: regressor.fit(X_train, Y_train)
```

```
Out[28]: LinearRegression()
```

Model Evaluation

```
In [29]: # prediction on training data
training_data_prediction = regressor.predict(X_train)
```

```
In [30]: # R squared value
r2_train = metrics.r2_score(Y_train, training_data_prediction)
```

```
print('R squared value : ', r2_train)
```

```
R squared value : 0.751505643411174
```

```
In [31]: # prediction on test data
test_data_prediction = regressor.predict(X_test)
```

```
In [32]: # R squared value
r2_test = metrics.r2_score(Y_test, test_data_prediction)
print('R squared value : ', r2_test)
```

```
R squared value : 0.7447273869684077
```

Building a Predictive System

```
In [33]: input_data = (31,1,25.74,0,1,0)
```

```
In [34]: # changing input data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)
```

```
In [35]: # reshape the array
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
```

```
In [36]: prediction = regressor.predict(input_data_reshaped)
print(prediction)
```

```
[3760.0805765]
```

```
C:\Users\Asus\anaconda\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names
, but LinearRegression was fitted with feature names
warnings.warn(
```

```
In [37]: print('The insurance cost is USD ', prediction[0])
```

```
The insurance cost is USD 3760.0805764960496
```

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js